



Campus Ride-Booking App Prototype – Technical Design and Architecture

1. Recommended Tech Stack for Rapid Prototyping

For a student project, use widely-adopted, easy-to-learn technologies that allow fast development. Common choices include:

- **Web Frontend:** React.js (JavaScript) or Vue.js for the web interface [1](#). React is very popular, and its ecosystem (e.g. React Router, Material-UI) speeds UI development.
- **Mobile Frontend:** React Native (JavaScript) or Flutter (Dart) for cross-platform mobile apps [1](#) [2](#). Both let you share code across iOS/Android. React Native ties directly to React web skills, while Flutter offers fast UI rendering.
- **Backend:** Node.js with Express (JavaScript) or Python with Flask/Django. Both support rapid API development [3](#) [4](#). Node.js (Express) is non-blocking and scales well with real-time data [5](#), whereas Python offers easy ML integration (TensorFlow/PyTorch). For prototyping, frameworks like Express or Flask are sufficient.
- **Database:** A managed cloud database like MongoDB Atlas (NoSQL) or PostgreSQL (SQL) [6](#) [7](#). MongoDB is schemaless and supports geo queries (via 2dsphere index) for finding nearby drivers. PostgreSQL with PostGIS can also handle location queries. A simple JSON document DB (Firestore or DynamoDB) is another option for quick setup.
- **Hosting/Backend DevOps:** Cloud platforms with free tiers simplify hosting. For example, Heroku or Vercel for deploying the frontend/backend, or AWS/GCP/Azure for more control [8](#) [9](#). Heroku's GitHub integration or Netlify/Vercel's automated deploys work well for student projects. Use GitHub for version control and GitHub Actions (or Jenkins) for CI/CD [10](#) [11](#). Docker can be used to containerize the backend for easier deployment [10](#).
- **Maps/Location:** Google Maps Platform (Maps JavaScript API and Geocoding API) for maps and distance calculations [12](#). Alternatives include Mapbox or OpenStreetMap. The apps can use the device's GPS (via HTML5 geolocation on web or **Geolocation** API in React Native) and then query the Maps API to find nearby drivers.

2. High-Level System Architecture

At a high level, the system consists of *clients*, a *backend service*, a *database*, and *third-party services*. The **clients** are the rider's web/mobile app and an admin panel. These communicate over HTTPS to the **Backend/API Server**, which implements application logic and exposes RESTful endpoints. The backend talks to the **Database** to store users, drivers, rides, and ratings. It also interacts with external services like the Maps API for location data and a SMS/calling API (e.g. Twilio) for driver contact. A simplified flow is:

- User App (Web/Mobile) ⇌ Backend API ⇌ Database
- Backend API ⇌ Google Maps/Geocoding API (for location queries)
- User App ⇌ (Phone/SMS) ⇌ Driver (direct contact)

Meanwhile, an **Admin Panel** (a web UI for operators) uses the same backend APIs to add or manage driver records and view analytics.

In words, the architecture comprises: - **Client Layer:** React web app and React Native (or Flutter) mobile app. They handle UI, user location permission, and calling backend APIs for login, search, and details.

- **API Layer:** A RESTful server (e.g. Node/Express or Python/Flask) handling authentication (JWT or OAuth2), business logic (finding nearby drivers, saving ratings), and calling external APIs.

- **Database Layer:** Cloud database (MongoDB/PostgreSQL) storing entities: users, drivers (with names, phone numbers, current coords), rides (if implemented), and ratings.

- **External Services:** Google Maps API for geocoding user location to lat/long and for computing distances¹²; Firebase Cloud Messaging or a push service for later notifications; Twilio (optional) for calling/SMS.

- **Admin Interface:** A separate React-based web dashboard for the “staff” role to manage driver listings and monitor usage. It uses the same backend but with admin-level APIs.

This modular setup allows each component to be developed and tested independently.

3. Component Architecture

- **Frontend (User App):** The user-facing apps (web and mobile) present maps, driver lists, and details.

Key components include login/register screens, a home map (showing user location), a list of *nearby drivers*, and a driver-detail page (with name, vehicle info, photo, and a “Call Driver” button). The web app could be a React single-page app, and the mobile app built with React Native or Flutter for iOS/Android. Both share the same backend APIs. In this MVP phase, the app simply displays the driver’s phone number for direct calling (using a `tel:` link).

- **Frontend (Admin Panel):** A protected web interface where an administrator can add or remove drivers, edit driver details (e.g. availability), and view simple stats. This can be built in React (possibly reusing components). It authenticates against the backend and calls admin-specific endpoints (e.g. `POST /admin/drivers`) to manage the database.

- **Backend API:** A server (Node.js/Express or Python/Flask) providing REST endpoints. It handles:

- **Authentication:** Endpoints like `POST /auth/register` and `POST /auth/login` to create users and issue JWTs¹³. (For simplicity, start with email/password or phone OTP login.)
- **Driver Search:** `GET /drivers?lat={x}&lng={y}` returns a list of drivers sorted by proximity (requires computing distances via a query or in-memory filtering).
- **Driver Details:** `GET /drivers/{id}` returns info (name, vehicle, contact, average rating).
- **Ratings:** `POST /drivers/{id}/ratings` to submit a user rating. (The server updates driver’s rating average.)
- **Admin APIs:** `POST /drivers` to add a new driver (admin only), `PUT /drivers/{id}` to update info, etc.
- **(Optional Rides):** If recording trip requests, `POST /rides` could save ride records after a user calls a driver. This is a future enhancement.

The backend also integrates Google Maps APIs (e.g. using the Maps JavaScript or Directions API) to handle geocoding and distance calculation on the server side if needed ¹².

- **Database:** A hosted DB instance (MongoDB Atlas or cloud SQL). Collections/tables might include *Users*, *Drivers*, *Rides*, *Ratings*. Each driver record stores a current or last-known location (latitude/longitude) and contact info. Use a geospatial index (e.g. MongoDB's 2dsphere index) to efficiently query nearby drivers. A separate **cache** (Redis) could be introduced if read performance becomes critical (e.g. caching the list of drivers).
- **Push/Notifications (Future):** In later phases, implement Firebase Cloud Messaging to send push alerts (e.g. "Your ride has arrived"). Initially this is optional; focus on the phone call flow first ¹⁴.
- **GPS Module (Future):** In a full version, driver and rider apps would periodically send GPS coordinates to the backend. This could be done via a background service in React Native or by polling the user's location in the web app. The backend would then stream updates (e.g. with WebSockets/Socket.IO) to allow a mini "live tracking" view. This is planned for a later phase and can be added without changing the core API design.
- **AI/Microservices (Future):** The architecture should allow plugging in separate AI services later. For example, you might have a standalone ML service (written in Python) that the backend can query for route optimization or demand prediction. These can be deployed as separate microservices or cloud functions. (See Section 6 for more on this.)

4. Deployment Stack & CI/CD

For student use, leverage PaaS and simple CI:

- **Version Control:** Host code on GitHub/GitLab. Use branches (`main` for production-ready code).
- **CI/CD Pipeline:** Set up GitHub Actions (free for public repos) or GitLab CI. On each push, run automatic checks (unit tests, linting). If tests pass, automatically deploy. Examples: Deploy the backend to Heroku using the [Heroku GitHub integration], or use a GitHub Action like `appleboy/heroku-action` ¹⁰. For the frontend, deploy to Netlify or Vercel via their Git connectors.
- **Containerization (Optional):** Use Docker to package the backend. For example, a Dockerfile can build the Node/Python server and its dependencies. Then you can push the image to Heroku (which supports container deploys) or to a cloud container registry. Docker guarantees consistency across dev and production. Kubernetes (e.g. Google Kubernetes Engine) is available for larger scale, but likely unnecessary for 500 users ¹⁰.
- **Automated Deployment:** A typical CI flow: developer pushes code → GitHub Actions runs tests → if green, Action deploys to Heroku/AWS (e.g. using `heroku container:push`) ¹⁰. Alternatively, connect the repo directly to Heroku and enable automatic deploys on `main`. Frontend code can auto-deploy via Netlify/Vercel on push.
- **Hosting:**

- **Backend:** Use Heroku (free tier dynos) or AWS Elastic Beanstalk/Lightsail. Heroku is easiest: just link to the repo and it builds Node/Python.
- **Database:** If using MongoDB, use MongoDB Atlas (free cluster). For PostgreSQL, use Heroku Postgres free tier or Google Cloud SQL. These managed DBs auto-scale modestly.
- **Frontend:** Host the web app on Netlify or Vercel (both have free tiers) for simple HTTPS serving of the React build. The mobile app would be deployed to simulators or device (not a "host").

By using Git-based deploys and Actions, we get continuous delivery with minimal setup. Tools like Docker and Kubernetes are available but optional for a student MVP ¹⁰.

5. API Endpoints (RESTful)

Key APIs for the MVP include:

• **Authentication:**

- `POST /auth/register` – Create a new user account (collect name, email/phone, password).
- `POST /auth/login` – Log in (return JWT token).

• **Driver Discovery:**

- `GET /drivers?lat={lat}&lng={lng}` – List drivers near the given location (optional radius).
The backend filters drivers by distance or bounding box using geospatial queries.
- `GET /drivers/{id}` – Get detailed info about one driver: name, vehicle, photo URL, phone number, average rating.

• **Driver Management (Admin):**

- `POST /drivers` – Add a new driver (admin only).
- `PUT /drivers/{id}` – Update driver info (admin only).
- `DELETE /drivers/{id}` – Remove a driver.

• **Ratings:**

- `POST /drivers/{id}/rating` – User submits a rating (1–5 stars) and optional comment.
Backend updates the driver's aggregate rating.
- `GET /drivers/{id}/rating` – (Optional) fetch all ratings or average rating for a driver.

• **User Profile (Optional):**

- `GET /users/me` – Return current user info (profile, ride history).
- `PUT /users/me` – Update profile (e.g. display name).

• **Other (Future):**

- Endpoints for handling ride requests/bookings (if implemented later), push notification tokens, or messaging.

All APIs should use JSON over HTTPS. Authentication tokens (JWT) are sent in the `Authorization` header. Follow REST conventions (use `GET` for reads, `POST/PUT` for changes). For example, the Uber API also uses a RESTful design with JSON responses ¹⁵.

6. Future AI Integration (Optional)

The architecture should be **extensible** to add AI/ML features. Possible AI-powered enhancements include:

- **Route Optimization & ETA:** Use an ML model to predict the fastest route or arrival time given current traffic. This could be a microservice that takes start/end coordinates and returns a route or ETA. (AI frameworks: TensorFlow or PyTorch model deployed on AWS SageMaker or Google Vertex AI ¹².)
- **Dynamic Pricing:** A demand-prediction model that adjusts fares during peak hours or events. This runs offline on historical ride data (user demand patterns) and exposes recommendations via an API.
- **Driver-Passenger Matching:** More intelligently match riders to drivers based on past ratings, preferences, and location, rather than just distance ¹⁶. An AI service could score possible matches (e.g. using machine learning on past ride data).
- **Chatbot Assistant:** Add an NLP chatbot (using a service like OpenAI GPT) in the app for booking help or FAQs. For instance, a rider could text “Book me a ride to the dorm” and an AI parses and acts on it. TechBuilder suggests using OpenAI APIs for such functions ¹².
- **Feedback Analysis:** Apply sentiment analysis on user feedback/comments to detect issues (long waits, rude drivers) and automatically flag problems.
- **Fraud/Safety Monitoring:** Use simple AI to detect anomalies (e.g. a driver’s GPS suddenly jumps far away, or multiple fake accounts) and alert admins ¹⁶.

To integrate AI, treat these as separate services: e.g. deploy a Python Flask/FastAPI service that loads the ML model and serves predictions. The main backend can call this service via a REST endpoint or AWS Lambda. Use cloud ML tools: TensorFlow/PyTorch for model building, and AWS SageMaker or Google Vertex AI for hosting models ¹². Ensure your data pipelines collect training data (ride logs, locations, ratings) so models can be improved.

Checklist for AI integration: Data collection → Model training → Deploy as service → Connect to backend API. For example, after launching, periodically train a route-ETA model on historical GPS and traffic data, then query it in real-time. As one source notes, startup teams can launch an MVP AI feature (like smarter route matching) in weeks, then iterate ¹⁶.

7. Security Considerations

Security is critical even for a prototype. Key practices include:

- **Transport Security:** Use HTTPS/TLS for all communications (frontend ↔ backend and backend ↔ external APIs) ¹⁷. Never send tokens or personal data over plain HTTP.

- **Authentication & Authorization:** Implement secure login with hashed passwords (bcrypt) and issue JWTs ¹³. Protect API endpoints by verifying JWTs. Use role-based checks so only admins can call admin APIs. Consider rate-limiting login attempts to prevent brute force.
- **Input Validation:** On the server, validate and sanitize all inputs to prevent SQL/NoSQL injection or XSS. For example, check that latitude/longitude query params are valid numbers. Always treat client data as untrusted.
- **Data Protection:** Store minimal PII. Encrypt sensitive fields in the database if needed. Do not hard-code API keys (Maps API key) in the client; keep them on the backend or in environment variables. For example, use environment variables or a secrets manager for any credentials.
- **Secure Storage (Client):** On the mobile app, store tokens securely (use SecureStore or Keychain). On web, use HTTP-only cookies or secure `localStorage` with short expiration.
- **Driver Contact Privacy:** If direct phone numbers are shown, consider future measures to mask them (e.g. Twilio's proxy numbers) to avoid misuse. For now, display numbers as provided, but caution around storing too many phone numbers.
- **Server Hardening:** Ensure the server and database require strong passwords. Use cloud-hosted DBs with SSL. Do not expose database ports to the internet; use backend authentication. Keep all dependencies up to date to patch known vulnerabilities.

As a guideline, follow OWASP recommendations (e.g. OWASP Mobile Cheat Sheet). An academic source emphasizes using *encryption in transit and at rest, access controls*, and multi-factor auth where possible. In summary: use HTTPS, JWT/OAuth2, bcrypt hashing, input validation and least-privilege ACLs for roles ¹⁷.

8. Scalability Plan (50→500 Users)

While 500 users is modest, it's wise to design for growth:

- **Load Testing:** Initially test the system with simulated load (e.g. JMeter or Locust) to identify bottlenecks. As one study notes, "simulate heavy traffic to evaluate scalability and performance under high load" ¹⁸. For example, verify the driver-search API still responds quickly with many users/driver entries.
- **Stateless Backend:** Build the API to be stateless (session info in JWT) so you can run multiple server instances behind a load balancer if needed. At ~500 users, a single dyno/container may suffice, but plan to scale out if traffic spikes.
- **Database Scaling:** Use a managed cloud DB that can scale vertically or with read replicas. Index frequently queried fields (e.g. location coordinates) for faster search. If using MongoDB, sharding or adding more Atlas cluster nodes can help. For SQL, ensure proper indexing and consider read replicas for reporting.
- **Caching:** Introduce Redis or in-memory caching for frequent queries (e.g. cache the list of available drivers for 30 seconds) to reduce DB load. Also use HTTP caching for static assets (CSS/JS) via a CDN.
- **Asynchronous Tasks:** Offload heavy work (like sending emails or SMS) to background jobs/queues (e.g. Redis queue with Bull or Sidekiq). This keeps the API responsive.
- **Container Orchestration (Future):** For larger scale, deploy with Kubernetes or a serverless framework. As noted, "cloud-native architectures (like Kubernetes or serverless) are often the best path forward" for flexible scaling ¹⁹. While overkill for 500 users, this mindset ensures you don't build hard limits.
- **Monitoring & Alerts:** Add basic monitoring (Prometheus/Grafana or even cloud watch metrics) to track server CPU, memory, and request rates. Set alerts if response time spikes.

- **Progressive Enhancements:** If one backend instance slows, you could split services (e.g. separate the rating service or AI service). However, it's probably fine to start monolithic and refactor later.

In practice, going from 50 to 500 users might only require increasing Heroku dyno size or DB tier, adding caching, and verifying indexes. At each step, re-run load tests and optimize slow endpoints (for example, ensure the driver-radius query uses a geospatial index). Following scalable practices (stateless APIs, containerization, managed DB) ensures smooth growth [10](#) [18](#).

Sources: Recommendations above combine industry best practices and recent guides for ride-hailing apps. For example, common stacks include React/Flutter on front-end and Node/Python on back-end [4](#) [3](#); Google Maps API is standard for location [12](#); and Firebase or AWS services (Heroku, GitHub Actions) simplify deployment [8](#) [10](#). Security and AI notes are drawn from expert blogs and case studies [17](#) [19](#). (All cited sources are publicly available.)

[1](#) [4](#) [7](#) [9](#) [11](#) [12](#) [19](#) How to Build a Game-Changing AI-Based Ride-Sharing App in 2025 - TechBuilder

<https://techbuilder.ai/how-to-build-a-game-changing-ai-based-ride-sharing-app-in-2025/>

[2](#) [3](#) [5](#) [6](#) [8](#) [10](#) [13](#) [14](#) [17](#) Best Tech Stack for Taxi Booking App Development Guide

<https://www.digitrix.com/blogs/best-tech-stack-for-taxi-app-development>

[15](#) Uber Developers

<https://developer.uber.com/docs/riders/references/api>

[16](#) AI-Powered Ride-Sharing App Guide | Build Smarter Mobility Solutions | Zestminds

<https://www.zestminds.com/blog/ai-powered-ride-sharing-app-guide/>

[18](#) ijnrd.org

<https://ijnrd.org/papers/IJNRD2505135.pdf>